Machine Learning for Radio Astronomy: Everything is binary classification if you phrase it right

Matthew Alger

Slides: http://www.mso.anu.edu.au/~alger/radio-lunch-may

Australian National University

We have too much data

- Surveys like VLA-FIRST generate more data than we can look at
- Surveys like ASKAP-EMU generate more data than we can *store*

Australian SKA Pathfinder.
Image: CSIRO

The Very Large Array.
Image: NRAO

Lots of data hold lots of astrophysics

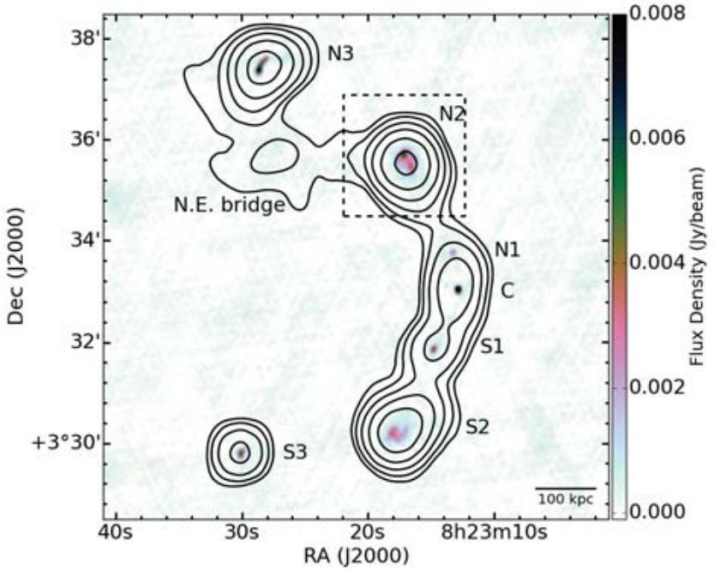

- Even 20-year old wide-area surveys like NVSS have lots of interesting astrophysics buried in them
- Much of this has come from manual inspection
- Plenty still to find

Bent, giant radio galaxy in NVSS/FIRST.
(~1 Mpc physical extent)
Image: Banfield+16

Strategies for data-at-scale

- Crowdsourcing (e.g. Radio Galaxy Zoo, Gravity Spy)
  - Fast: RGZ has classified 75000 galaxies in just 3 years
  - Serendipitous: Citizen scientists are endlessly curious
  - Noisy: Non-experts are not experts
- Asking students very nicely to look at all the data
  - Slower: Students are slow and grumpy
  - Opportunity cost: More fruitful things to do
  - Incomplete: We can't see *all* the data, so we miss things
- Machine learning
  - Fast: Computers are well-known to be quite speedy
  - Hard to interpret: Much state-of-the-art ML research is black magic
  - Unclear how to develop: Given a problem, how do we make ML work for it?

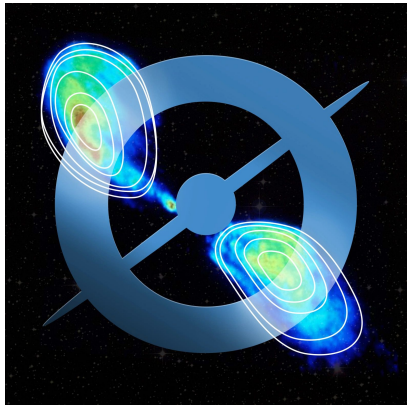

Obligatory xkcd

- “Machine learning has become alchemy.”
  — Ali Rahimi,
  NIPS 2017 Test-of-Time
- “[People] underestimate how much can be achieved with relatively crude systems”
  — François Chollet
- For useful science results, we need to understand what our methods are doing
  - Different to understanding how they are doing it!

Image: xkcd

Machine learning in radio

- Visualisation (Polsterer+15)
- Source classification (Aniyan+17)
- Component classification (Lukic+18)
- Host galaxy cross-identification (Alger+prep)
- Source identification (Wu+prep)

Host Galaxy Cross-Identification

- Problem:
  - Match radio emission to its host galaxy at other wavelengths
- Hard:
  - Radio emission can be extended at scales of tens of arcminutes
  - Often no clear relationship between radio emission and host galaxy

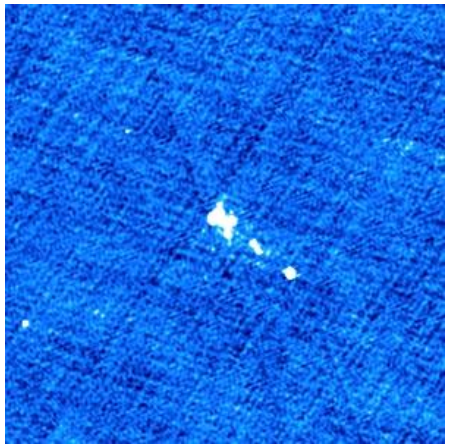

FIRSTJ023838.0+023450 at 1.4 GHz.
Image: FIRST

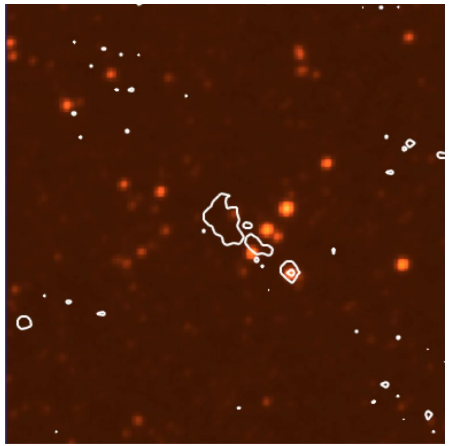

FIRSTJ023838.0+023450 in infrared.
Image: WISE

Machine learning can only answer some questions

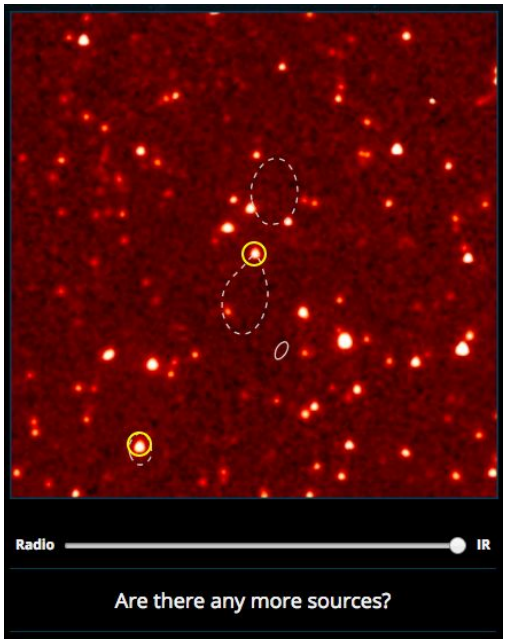

How do you turn an astrophysics question like “Where’s the host galaxy?” into a machine learning question like “Is this a 1 or a 0?”

$$f : something \rightarrow \{1, 0\}$$

$$host : radio\ emission \rightarrow IR\ host\ ?$$

Learning to cross-identify radio emission

f : IR galaxy → {host, not host}

- Radio Galaxy Zoo = hosts
- Can use basic (and simple!) machine learning techniques to train and test models

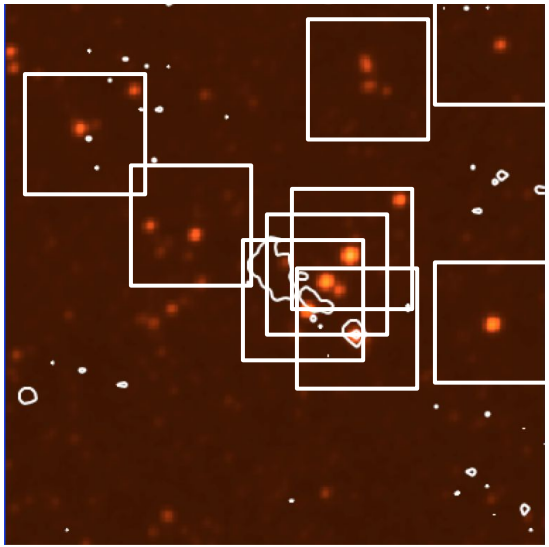

Candidate host galaxies.
Image: FIRST/WISE

Cross-identification with binary classification

(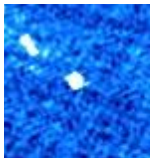 , 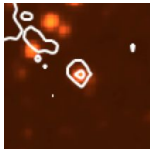) ⟶ 1

(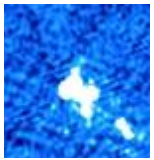 , ) ⟶ 0

Representation of galaxy

Whether galaxy has an AGN

Understanding the link between ML and physics

- Since applying ML requires you force your problem into an ML framework, performance measures become confusing
  - e.g. classification accuracy to cross-identification accuracy
  - Uncertainties
- Evaluation and baselines important but underrated